

03_01_Synchronization

1. Background

1.1. Neccissity

1.2. Achievement

2. Happened-before Relation (Causality Relation)

2.1. Happened-before Relation (Causality Relation)

2.2. Concurrency Relation:

3. Logical Clocks

3.1. Requirements

Consistent

Characterizes

3.2. Scalar Clocks (Lamport Clocks)

Problem

3.3. Vector Clocks

Vector Clocks

Prepare

Implementation Rules

Meaning

Characterize Analysis

3.4. Characterizing the HB relation

Theorem: (See Slides)

1. Background

1.1. Neccissity

Computers may need a notion of time

- for **assigning timestamps** to events
- for **comparing timestamps** to order events

1.2. Achievement

In asynchronous distributed systems, synchronization has to be **achieved through message**

2. Happened-before Relation (Causality Relation)

2.1. Happened-before Relation (Causality Relation)

The **happened-before relation** (or **causality relation** \longrightarrow) on E is the the smallest relation satisfying

- **local order:** if a, b in E_i , and a occurs before b , then $a \longrightarrow b$
- **message exchange:** if a in E_i is the event of sending a message m and b in P_j is the event of receiving m , then $a \longrightarrow b$
- **transitivity:** if $a \longrightarrow b$ and $b \longrightarrow c$, then $a \longrightarrow c$

2.2. Concurrency Relation:

if neither $a \longrightarrow b$ nor $b \longrightarrow a$ holds, then a and b are **concurrent** (notation: $a \parallel b$)

3. Logical Clocks

3.1. Requirements

Timestamps have to obey the HB relation

Consistent

A logical clock is a function $C: E \longrightarrow S$ which is **consistent** with the HB relation:

if $a \rightarrow b$, then $C(a) < C(b)$

$$\text{if } a \rightarrow b, \text{ then } C(a) < C(b)$$

Characterizes

A logical clock **characterizes** the HB relation if:


$$C(a) < C(b) \text{ iff } a \rightarrow b$$

3.2. Scalar Clocks (Lamport Clocks)

- Assume an **integer counter** C_i in process P_i
- Initialization of the C_i is irrelevant

Two **implementation rules**:


1. If a in E_i and a is **not a message-receive event**, then
 - P_i first **increments C_i** (e.g., by 1)
 - and then **sets $C(a)$ to the new value of C_i**
2. If a is the event in P_i of sending message m and b is the event in P_j of receiving m , then
 - P_i sends $C(a)$ along with m
 - P_j assigns C_j the value **$\max(C_j+1, C(a)+1)$**
 - P_j sets $C(b)$ to the new value of C_j



Problem

A scalar clock **cannot characterize** the HB relation:

- take a set of processes with **only internal events** (no messages)
- any pair of events in any two different processes are **concurrent**, but will in general not have equal clock values



3.3. Vector Clocks

Vector Clocks

- **Comparison** of two vectors v, w of length k :

$$\begin{aligned}
 v = w & \quad \text{iff } v[i] = w[i], \quad i = 1, 2, \dots, k \\
 v \leq w & \quad \text{iff } v[i] \leq w[i], \quad i = 1, 2, \dots, k \\
 v < w & \quad \text{iff } v[i] \leq w[i], \quad i = 1, 2, \dots, k, \text{ and } v \neq w \\
 v \geq w & \quad \text{iff } v[i] \geq w[i], \quad i = 1, 2, \dots, k \\
 v > w & \quad \text{iff } v[i] \geq w[i], \quad i = 1, 2, \dots, k, \text{ and } v \neq w
 \end{aligned}$$

- Component-wise **maximum**

$$\max(v, w)[i] = \max(v[i], w[i])$$

- **Unit vector** in dimension i :

$$e_i = (0, 0, \dots, 0, 1, 0, \dots, 0)$$

Prepare

Assume a **vector of integers** V_i in process P_i

Initialization: $V_i = (0, 0, \dots, 0)$

Implementation Rules

1. If a in E_i and a **is not a message-receive event**, then
 - P_i first increments $V_i[i]$ by 1 /* count local events */
 - and then sets $V(a)$ to the new value of V_i
2. If a is the event in P_i of sending message m and b is the event in P_j of receiving m , then /* next event in P_j */
 - P_i sends $V(a)$ along with m
 - P_j assigns V_j the value $\max(V_j + e_j, V(a))$
 - P_j sets $V(b)$ to the new value of V_j

Meaning

- In $V_i[i]$, P_i simply numbers its local events
- The meaning of $V_i[j]$ with $i \neq j$ is: the number of the last event in P_j that P_i “**knows about**,” either directly or indirectly
- Two events a and b are **concurrent** if there two indices i and j such that
 - $V(a)[i] > V(b)[i]$
 - $V(a)[j] < V(b)[j]$
 - Above means timestamps of a and b are incomparable: at event a there is more “knowledge” about process P_i and at event b there is more “knowledge” about P_j

Characterize Analysis

- if $a \rightarrow b$, then $V(a) < V(b)$
- if $V(a) < V(b)$:
 - suppose a occurs in P_i and b in P_j
 - because $V(b)[i] \geq V(a)[i]$, at event b there is **at least the same amount of knowledge** about P_i as at event a
 - this can only be caused by a chain of events from a to b , so $a \rightarrow b$

3.4. Characterizing the HB relation

Theorem: (See Slides)

For a **k-dimensional vector clock** to characterize the HB relation in a system with n processes, we need $k \geq n$, that is, **size of vector clock is at least equal to the number of processes**