

03_03_Message Ordering

1. Introduction

1.2 Causal Order

1.2. Total Order

1.3. Relations

1.4. Basic Idea of Message Ordering Algorithm

2. Algorithms for Causal Order

Assumption

2.1. Causal ordering for broadcast messages (Birman-Schiper-Stephenson algorithm)

Main Idea

Process

Understanding

Implementation

Correctness

2.2. Causal ordering for point-to-point messages (Schiper-Eggli-Sandoz algorithm)

Prepare (Structure)

Main Process

Understanding

Implementation

3. Algorithms for Total Ordering

3.1. Simple Solution

3.2. Total ordering for broadcast messages

Main Idea

Delivery Condition

Correctness

1. Introduction

For a message m , the set of destinations is denoted by $Dest(m)$.

- **Point-to-point:** $|Dest(m)| = 1$
- **Multicast:** $|Dest(m)| > 1$
- **Broadcast:** $|Dest(m)| = \#processes$

The event of multicasting a message m is denoted by $\mathbf{m}(m)$


The event of **delivering** a message m to process P_i is denoted by $d_i(m)$

1.2 Causal Order

Message order is causal when for every two messages m_1 and m_2 :

- if $m(m_1) \rightarrow m(m_2)$,
- then $di(m_1) \rightarrow di(m_2)$ for **all i in both Dest(m_1) and Dest(m_2)**

Example:



The contents of m_2 can depend on the contents of m_1 , so the delivery of m_2 in P_3 has to be postponed until after m_1 has been delivered

1.2. Total Order


Message order is total when all processes receive all messages in the same order (if broadcast)

1.3. Relations

- Total order **does not imply** causal order
- Causal order **does not imply** total order

1.4. Basic Idea of Message Ordering Algorithm

For message ordering, assume an **additional process** P' for every process P in the system:



- when a message arrives, P' checks whether the message can be delivered to P according to the required order
- if not, P stores the message in a buffer, and re-checks when messages that arrive later have been delivered
- So **reception and delivery** of a message are **separate**

2. Algorithms for Causal Order

Assumption

FIFO Channels

2.1. Causal ordering for broadcast messages (Birman-Schiper-Stephenson algorithm)

Main Idea

- every process **numbers** its own broadcast messages

- a receiving process **checks** whether it is the broadcast message it expects from the sending process
- processes **transfer knowledge** about broadcast messages they have received in the timestamps of the messages they send
- a receiving process **checks** whether it has not missed messages

Process

- A message m is accompanied by the value V_m of the local vector clock when it is sent
- The condition for **delivery** of a message m in P_i from P_j is

»

Understanding

The delivery condition means:

- the message is the next one expected from P_j (equality in $D_j(m)$)
- with respect to all other processes, P_i is **at least as up to date as** P_j was when it sent the message

Implementation

I. Broadcasting a message

```
V := V + e_i      /* first increment local clock */
broadcast(m, V)
```

II. Receiving a message from P_j

```
upon receipt of (m, V_m) do
  if D_j(m) then
    deliver(m)
    while ( |{(m, k, V_m) in B | D_k(m)}| > 0 ) do
      deliver such a message m
  else add (m, j, V_m) to B
```

check for other
messages that
can be delivered

III. Delivering a message from P_j function deliver(m)

```
deliver(m) to P_i
V := V + e_j      /* only modification in component j */
remove(m, V_m) from B /* in the other components the local vector clock */
/* is at least equal to the message clock anyway */
```


Correctness

- **Correctness of the Berman-Schiper-Stephenson algorithm:**

- for broadcasts sent by the same process, trivial

- now suppose

- m_1 sent by P_j
 - m_2 sent by P_k
 - both received by P_i
 - $m(m_1) \rightarrow m(m_2)$



- so $V(m_2)[j] \geq V(m_1)[j]$ **(**)** **(*)** **(**)**
 - condition for delivery of m_2 in P_i : $V_i[j] \geq V(m_2)[j]$ ($\geq V(m_1)[j]$)
 - but $V_i[j]$ is only ever modified (incremented by 1) by receiving broadcasts from P_j , so P_i must first receive m_1

2.2. Causal ordering for point-to-point messages (Schiper-Eggli-Sandoz algorithm)

Prepare (Structure)

- Use **vector clocks** (in the ordinary way), all initialized to all **zeroes**
- Every process maintains a **local buffer S** of ordered pairs each consisting of a **process id** and a **timestamp**
- A **message** sent by P_i is accompanied by the **complete current contents of the local buffer S_i**

Main Process

- Condition $D_i(m)$ for delivery of message m with accompanying buffer S_m in P_i :
 - there does not exist (i, V) in S_m
 - or there does exist (i, V) in S_m and $V \leq V_i$
- When a message is delivered in P_i , the **knowledge carried by the message** and the **knowledge available in P_i** about all other processes **are merged**
- Time to update local buffer:
 - when sending
 - when delivering

Understanding

- Meaning of this pair (P_j, V_j) in P_i :
 - the most recent knowledge in P_i about what P_j should know

- so can be used to tell P_j what it should know
- V_j is at least as large as the timestamp of the last message from P_i to P_j
- The reason message accompanied by buffer; indicating to the receiving process what P_i thinks others should know
- $D_i(m)$
 - either m carries no knowledge at all about what P_i should have received
 - or m does carry such information, but P_i is sufficiently up to date

Implementation

I. Sending a message to P_j :

send(m, S, V) to P_j

insert(j, V) into S /* delete any old element for P_j */

II. Receiving a message

same as for previous algorithm

III. Delivering a message in P_i :

deliver(m) to P_i

for all ((j, V')) in S_m) do

if (there exists (j, V'') in S) then

remove (j, V'') from S

$V'' := \max(V', V'')$

insert(j, V'') into S

else insert(j, V') into S

merge local buffer with
buffer in message

3. Algorithms for Total Ordering

3.1. Simple Solution

have a **special process** P_0 (a **sequencer**)

- when a process wants to do a broadcast, it sends the message to P_0
- P_0 numbers all messages sequentially, and then broadcasts them to all processes
- P_0 keeps a history of messages so that if a process misses a message, it can ask for a resend

3.2. Total ordering for broadcast messages

Main Idea

- use **scalar clocks** (with process ids as **tie breakers**)
- all messages carry a **timestamp**

$(ts1, pid1) < (ts2, pid2)$
if $ts1 < ts2$ **or**
if $ts1 = ts2$ **and** $pid1 < pid2$

- every process maintains an ordered **message queue**
- **all** processes **acknowledge all** messages to **all** processes (include sender and itself)

Delivery Condition

A message can be delivered in a process when:

1. it is at **the head of** the local message queue (it is **the oldest message** the process knows about)
2. the process has **received an ack** for that message **from every process** (so no older message will arrive)

Correctness

because of acks and FIFO, all previous messages are forced out to the process