

# 03\_06\_Deadlock Detection

## 1. Definition

### 1.1. General Definition

### 1.2. LTS perspective

#### 1.2.1. Conditions of Deadlock

## 2. Background

### 2.1. Why Deadlock Detection

#### 2.1.1. Three ways to deal with Deadlock

### 2.2. Types of Request

## 3. Wait-For Graph

### 3.1. Resource Mode

#### Representation

#### Deadlock

### 3.2. Communication Model

#### Representation

#### Deadlock

## 4. Algorithms-1

### 4.1. Candy, Misra and Haas AND model (Resource model)

#### Basic Idea

#### Understanding and correctness

#### Process

#### Implementation

### 4.2. Candy, Misra and Haas OR model (Communication Model)

#### Main Idea

#### Understanding

#### Process

#### Preparation (Structure)

#### Implementation

## 5. Bracha and Toueg

### 5.1. Bracha and Toueg with instaneous message

#### Main Idea

#### Model

#### Preparation (Structure)

#### Process

#### Implementation

#### Correctness

#### understanding

### 5.2. Bracha and Toueg without instaneous message

#### Assumption

#### Main Idea

#### Remark

## 1. Definition

## 1.1. General Definition

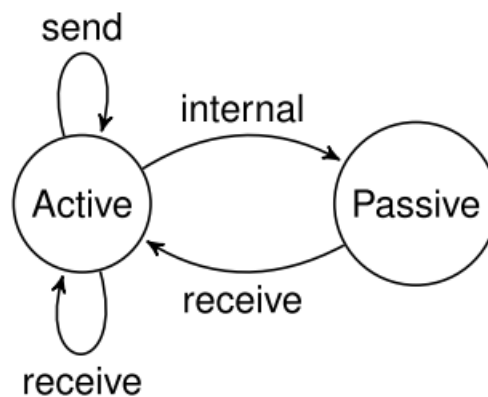
In a transition system a configuration is a **deadlocked or terminating** if there is **no outgoing transition** from that configuration

1. **terminating configuration** is reached when the algorithm **successfully completes** its task

## 1.2. LTS perspective

A node in isolation, can remain in one of the two states:

1. A process is active, if it has something to execute.
2. Otherwise, it is passive. A passive process will not necessarily always remain passive.



LTS\_perspective

### 1.2.1. Conditions of Deadlock

1. every process is in a passive state
2. all channels are empty

## 2. Background

If the initiator does not hear from any one for some time, how can it determine whether:

1. the computation is still making progress;
2. it is deadlocks, or
3. it has terminated?

### 2.1. Why Deadlock Detection

If deadlock is detected one can use a **rollback mechanism**

#### 2.1.1. Three ways to deal with Deadlock

## 2.2. Types of Request


- **OR request:** one positive reply is sufficient
- **AND request:** all processes must reply positively
- **N-out-of-M request:** (at least) N positive replies needed

## 3. Wait-For Graph


### 3.1. Resource Mode

#### Representation

- arrow from process P to resource Q (process  $R_Q$ ) if P has a request for Q that has not been granted
  - arrow from resource Q to process P if P has acquired Q



been granted:



been released:




### 3.2. Communication Model

#### Representation

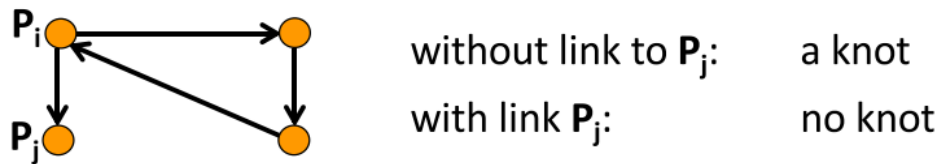
- arrows from process P to all processes from which P expects a single message
- So every blocked process  $P_i$  has a **dependent set**  $D_i$  of processes on which it waits

- A process turns active when it receives a message **from any** process in its dependent set



## Deadlock

- a cycle does not necessarily indicate a deadlock
- There is a deadlock iff **there is a knot** in the WFG
- A **knot**: a set  $S$  of nodes with:
  - a path from every node in  $S$  to every other node in  $S$
  - no edge from a node in  $S$  to a node not in  $S$



## 4. Algorithms-1

### 4.1. Candy, Misra and Haas AND model (Resource model)

#### Basic Idea

- when a process **suspects deadlock**, it sends a **special message** to all processes it is waiting for
- those processes **recursively propagate** these messages
- Deadlock if a process receives a message initially from itself


#### Understanding and correctness

- A process waits for all resources it has requested

- Process  $P_i$  is dependent on process  $P_j$  if **there is a path** in the WFG from  $P_i$  to  $P_j$

Processes maintain a boolean array:

- $\text{dep}_i(j)$  is true if and only if  $P_i$  knows that  $P_j$  is dependent on it



If  $\text{dep}_i(i)$  is true,  $P_i$  is **deadlocked**

Attention: This knowing is dynamically, that means the link can always exist, but if  $P_i$  has executed, then the value will be false (not "bottleneck")

### Process

- A process that suspects it is deadlocked initiates a **probe** by sending out probe messages to the processes it is waiting for
- Processes receiving a probe message **propagate it further** to the processes they are waiting for
- A process is deadlocked when it receives a probe message corresponding to a probe initiated by itself
- A probe message initiated by process  $P_i$  and sent by process  $P_j$  to process  $P_k$  has parameters  $(i, j, k)$
- $P_j$  sends probe  $(i, j, k)$  if
  - $P_j$  is blocked
  - $P_j$  is waiting for  $P_k$
  - $P_j$  knows that  $P_i$  is dependent on it (except when  $i=j$ , when the probe is initiated)

### Implementation

## I. Initiating a probe


```
for all (  $P_j$  for which  $P_i$  is waiting ) do  
    send( probe( $i, i, j$ ) )           /* initiator, source, destination */
```

## II. Process $P_i$ becomes executing

```
for all  $j$  do  $dep_i(j) := \text{false}$            /* reset dependencies */
```

## III. Receiving a probe message by $P_i$


```
upon receipt of ( probe( $j, k, i$ ) ) do  
    if ( (  $P_i$  is blocked ) and (  $dep_i(j) = \text{false}$  ) and  
        (  $P_i$  has not replied to all requests of  $P_k$  ) ) then  
         $dep_i(j) := \text{true}$            /* record dependency */
```



```
    if (  $i = j$  ) then  $P_j$  is deadlocked /* receive own probe */
```

```
    else
```

```
        for all (  $P_l$  for which  $P_i$  is waiting ) do  
            send( probe( $j, i, l$ ) )           /* propagate probe */
```



Note:

if  $P_i$  has replied to requests of  $P_k$ , and still receive probe ( $j, k, i$ ), that means the waiting status of  $P_j$  is not caused by  $P_i$ , and because  $dep()$  represents the current dependent situation

## 4.2. Candy, Misra and Haas OR model (Communication Model)

### Main Idea

- When a process **suspects deadlock**, it sends a **special message** (a "query") to all processes it is waiting for
- those processes **recursively propagate these messages**
- when a process has already received the message, **it sends back a reply immediately**
- when a process has exhausted all its links, **it sends back a reply**
- **deadlock** if the process receives a reply from all processes it sent a query to

## Understanding

A knot is a "closed universe", we can always find a tree and so for all queries a reply will be received. If there is a link go out and break the universe, from the implementation, this query will not return a **REPLY** because in segment 2, the first block will be executed, no reply will be send

## Process

- A process that suspects deadlock initiates a **query** in the system
- Queries of the same process **are numbered consecutively**
- **Reply** messages may be sent in reply to query messages
- **Parameters** of queries and replies are  $(i, m, j, k)$ :
  - $i$  is the id of the initiator of the query
  - $m$  is the query sequence number of  $P_i$
  - the message is sent by  $P_j$  to  $P_k$
- Structure of the algorithm:
  - wave of query messages along the "dependent-set" links
  - wave of reply messages back along these links (perhaps)
- A process that hears about a certain query for the first time, is "**engaged**" and **propagates** the query further
- A process that receives a query while it has **already been engaged**, sends a **reply back** (avoid cycles, create a tree)
- So a tree with **engager links** is created **until the dependent-set links** have been exhausted
- A process is **deadlocked** if and only if for every QUERY message it has received a corresponding REPLY message

## Preparation (Structure)

Data structures in every process  $P_i$ :

- **latest(j)**: the **largest query sequence number** of process  $P_j$  it knows about
- **engager(j)**: the id of the process that notifies it about **latest(j)**
- **num(j)**: ( # of query messages sent ) **minus**  
( # of reply messages received ) for the current query of process  $P_j$
- **wait(j)**: is true if it has been idle ever since it was notified of **latest(j)**

## Implementation

### I. Initiating a query (on suspicion of deadlock)

```
latest(i) = latest(i)+1           /* increase query number */
wait(i) = true
for all  $j \in D_i$  do send( query(i,latest(i),i,j) ) /* send out query */
num(i) =  $|D_i|$                    /* number of replies needed */
```

### II. Receiving a query message by $P_i$ from $P_k$

```
upon receipt of query(j,m,k,i) do /* process  $P_j$  is initiator */
  if ( m > latest(j) ) then        /* new query number */
    latest(j) = m                 /* record it */
    engager(j) = k                /* and record id of "engager" */
    wait(j) = true
    for all  $l \in D_i$  do           /* propagate query */
      send( query(j,m,i,l) )      /* to own dependent set */
      num(j) =  $|D_i|$ 
    else if ( (m=latest(j)) and wait(j) ) then /* if engaged before */
      send( reply(j,m,i,k) )      /* send reply back */
```

A message of the same initiator and the same query number may come back, so the first if is designed



### III. Receiving a reply message by $P_i$ from $P_k$

```
upon receipt of reply(j,m,k,i) do      /* process  $P_j$  is initiator */
if ( (m = latest(j)) and wait(j) ) then /* same query, still in wait */
    num(j) := num(j) - 1
    if ( num(j) = 0 ) then                /* all replies received */
        if ( j = i ) then                /* if initiator, */
             $P_i$  is deadlocked            /* then deadlocked, */
        else                             /* else */
            l := engager(j)              /* send reply back */
            send(reply(j,m,i,l))         /* only to engager */
```

### IV. On becoming executing

```
for all j do wait(j) := false          /* stop all messaging */
```

If someone in the "chain" is executing, no message will propagate by it.

## 5. Bracha and Toueg

### 5.1. Bracha and Toueg with instantaneous message

#### Main Idea

- Any process (the initiator) that **suspects a deadlock** of the basic computation **can initiate** the algorithm
- Two phases:
  - notify**: notify nodes that the algorithm has started (create a spanning tree)
  - simulate**: simulate the granting of resources
    - may unblock processes, which then also grant resources
    - simulates a continuation of the basic computation
- When the **initiator remains blocked** after the execution of the algorithm, it is **deadlocked**

#### Model

- Processes have two states:
  - active**: not waiting for another process
  - blocked**: waiting for a request to be satisfied (process cannot do any further requests)
- Channels are FIFO

- A **transformation** on the WFG corresponds to an action in a single process (e.g., doing a request adds edges to G)
- Let  $\sigma$  be a **schedule** (sequence of transformations)
- A node  $v$  is **active** if  $nv=0$
- **Transformations of a configuration G: 1**
  - **adding k outgoing edges** to  $v$  and setting  $nv=r$ 
    - result of doing an r-out-of-k **REQUEST**
  - **deleting an edge (v,u)** and decreasing  $nv$  by 1. If then  $nv=0$ , all remaining outgoing edges from  $v$  are **deleted** •
    - result of **sending a REPLY** to  $v$

### Preparation (Structure)

- Variables in a node:
  - $nv$ : number of REPLIES  $v$  still needs
  - $OUT_v$ : set of nodes  $u$  such that  $(v,u)$  in WFG
  - $IN_v$ : set of nodes  $u$  such that  $(u,v)$  in WFG
- Messages:
  - NOTIFY: wave of outgoing messages at the start involving other processes in the algorithm
  - DONE: finish of **NOTIFY** process (also end of grant if able)
  - GRANT: grant of resources, means sender is not blocked
  - ACK: reply to NOTIFY

### Process

### Implementation

**I. Initialization in node v:**

```
OUT := {u | (v,u) in E}      /* edges from v in WFG */
IN := {u | (u,v) in E}      /* edges towards v in WFG */
notified := false           /* not yet participated in notify */
free := false               /* deadlocked or not */
#granted := 0               /* number of requests granted */
```

**II. notify() procedure:**

```
notified := true             /* start of algorithm: */
for all (w in OUT) send(w,NOTIFY) /* initiator executes notify() */
if (n=0) then grant()        /* spread NOTIFYS */
for all (w in OUT) await(w,DONE) /* if active, simulate REPLY */
/* end of algorithm: */
/* initiator receives all DONEs */
```



**III. Receiving a NOTIFY in v from u:**

```
if (not notified) then notify() /* NOTIFYS spread outwards */
send(u,DONE)                    /* DONEs go back */
```

**IV. Procedure grant():**

```
free := true                   /* execute grant() only once */
for all (w in IN) send(w,GRANT) /* so simulate REPLY */
for all (w in IN) await(w,ACK) /* and wait for ACK */
```

**V. Receiving a GRANT by v from u:**

```
#granted := #granted+1
if ( (not free) and (#granted ≥ n) ) then grant() /* apply grant() */
send(u,ACK)                                     /* only once */
```

**Attention:**

- wait operations are **non-blocking**

- the initiator is **not deadlocked** if and only if at the end of the algorithm, **free=true** in the initiator
- sending **GRANT**s starts at nodes with **n=0**

## Correctness

See Slides

## understanding

- only a node with  $n=0$  can start **GRANT**
- DONE** message means the over of the algorithm
- GRANT** messages are used to **simulate REPLY** messages.

## 5.2. Bracha and Toueg without instaneous message

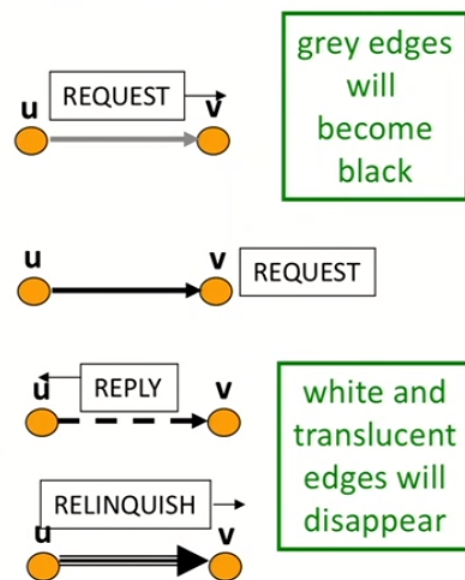
### Assumption

There maybe message in transit

### Main Idea

- 4 kind of Color of edge

- Color of edge  $(u,v)$ :
  - grey:
    - $u$  has sent a REQUEST to  $v$
    - $v$  has not received this
    - $u$  has not sent a RELINQUISH
  - black:
    - $v$  has received a REQUEST from  $u$
    - $v$  has not sent a REPLY to  $u$
    - $u$  has not sent a RELINQUISH
  - white:
    - $v$  has sent a REPLY to  $u$
    - $u$  has not received this REPLY
    - $u$  has not sent RELINQUISH
  - translucent:
    - $u$  has sent a RELINQUISH to  $v$
    - $v$  has not yet received it



- grey edges will become black
- white and translucent edges will disappear

- Transmation

- Transformations of a configuration **G**:


1. adding **k grey** outgoing edges to **v** and setting  $n_v=r$ 
  - doing an **r-out-of-k REQUEST**
2. changing a **grey** edge into a **black** edge
  - receiving a **REQUEST**
3. changing a **black** edge **(u,v)** into a **white** edge
  - sending a **REPLY** to **u**
4. removing a **white** edge **(u,v)** and decreasing  $n_u$ . If  $n_u=0$ , then all outgoing edges of **u** are made **translucent**
  - receiving a **REPLY** and possibly sending **RELINQUISHes**
5. removing a **translucent** edge
  - receiving a **RELINQUISH**

- Color Information

- Every process sends a **COLOR message** to all processes in its IN and OUT sets, making those processes aware of their membership of these sets.

1. **u** in **IN<sub>v</sub>** and **v** in **OUT<sub>u</sub>**:

- **v** has received **u's REQUEST**
- **(u,v)** is **black**




2. **u** in **IN<sub>v</sub>** and **v** not in **OUT<sub>u</sub>**:

- **v** does not yet know that **u's** request has been satisfied
- **(u,v)** is **translucent**

### 3. $u$ not in $IN_v$ and $v$ in $OUT_u$ :

- REQUEST still on its way ( $u$  not yet in  $IN_v$ ), or
- REPLY has been sent but is still on its way ( $u$  already removed from in  $IN_v$ )
- $(u,v)$  is **grey** or **white** (is all that's needed)



- Objective
  - Assume (**optimistically**) that:
    - the grey, white, and translucent edges are **non-existent**
    - so all REQUESTs corresponding to **grey and white edges** will be granted
  - Let **#greywhite** be the total number of outgoing grey and white edges
  - Then: a node  $v$  is **active** if  $n_v \leq \text{\#greywhite}_v$

#### Remark

- If a deadlock is detected then there was indeed one
- There may be a deadlock without detecting it at current time