

RL using Function Approximation_ For continuous space

RL using Function Approximation_ For continuous space

1. Q-function Approximation

1.1. Background & Basic Idea

1.2. Method 1: use linearly parametrized approximation

Q-function Approximation

Action Space Discretization

1.3. Fuzzy Q-iteration

Model

Policy

Fuzzy Q-iteration

2. Actor-critic methods

2.1. Structure

2.2. Update Critic: Value Estimation

2.3. Update Actor: Policy Update

3. One Really Classical Example

1. Q-function Approximation

1.1. Background & Basic Idea

In real-life control, X , U continuous

- approximate Q-function \hat{Q} must be used

1.2. Method 1: use linearly parametrized approximation

Q-function Approximation


- use **Basis Function** to approximate Q-function in a continuous space

$$\hat{Q} = \sum_{i=1}^N \theta_i \phi_i(x, u) \quad (1)$$

$$\phi_i(x, u) : X \times U \mapsto \mathbb{R}$$

usually normalized: $\sum_i \phi_i(x) = 1$

- E.g., fuzzy approximation, RBF network approximation



- Policy is greedy in \hat{Q} , computed on demand for given x :


$$\pi(x) = \arg \max_u \hat{Q}(x, u)$$

- Approximator must ensure efficient arg max solution

Action Space Discretization

- Choose M discrete actions $u_1, \dots, u_M \in U$
- Solve “arg max” by explicit enumeration
- In a control problem, we always make the grid more detail around the attractor, for example, 0 in the graph

Example: grid discretization



1.3.Fuzzy Q-iteration

Model

Given:

- N basis functions ϕ_1, \dots, ϕ_N
- M discrete actions u_1, \dots, u_M

Store

- N × M matrix of parameters θ (one for each pair basis function–discrete action)
 - same row: same ϕ
 - same column: same action

$$\hat{Q}^\theta(x, u_j) = \sum_{i=1}^N \phi_i(x) \theta_{i,j} = [\phi_1(x) \dots \phi_N(x)] \begin{bmatrix} \theta_{1,j} \\ \vdots \\ \theta_{N,j} \end{bmatrix} \quad (2)$$

Policy


$$\hat{\pi}^*(x) = \arg \max_{u_j, j=1, \dots, M} \hat{Q}^{\theta^*}(x, u_j) \quad (3)$$

(θ^* = converged parameter matrix)

Fuzzy Q-iteration

Fuzzy Q-iteration

```
repeat at each iteration  $\ell$ 
  for all cores  $x_i$ , discrete actions  $u_j$  do
     $\theta_{\ell+1,i,j} = \rho(x_i, u_j) + \gamma \max_{j'} \hat{Q}^{\theta_\ell}(f(x_i, u_j), u_{j'})$ 
  end for
until convergence
```


 weighted sum

2. Actor-critic methods

2.1. Structure

Explicitly separated value function and policy

- Actor = control policy $\pi(x)$
- Critic = state value function $V(x)$



Actor-critic

for every trial **do**

 initialize x_0 , choose initial action $u_0 = \tilde{u}_0$

repeat for each step k

 apply u_k , measure x_{k+1} , receive r_{k+1}

 choose **next** action $u_{k+1} = \hat{\pi}(x_{k+1}, \varphi_k) + \tilde{u}_{k+1}$

$\Delta_k = r_{k+1} + \hat{V}(x_{k+1}, \theta_k) - \hat{V}(x_k, \theta_k)$

$\theta_{k+1} = \theta_k + \alpha_c \Delta_k \left. \frac{\partial \hat{V}(x, \theta)}{\partial \theta} \right|_{\substack{x=x_k \\ \theta=\theta_k}}$

$\varphi_{k+1} = \varphi_k + \alpha_a \Delta_k \tilde{u}_k \left. \frac{\partial \hat{\pi}(x, \varphi)}{\partial \varphi} \right|_{\substack{x=x_k \\ \varphi=\varphi_k}}$

until terminal state

end for

When facing continuity:

- Actor parameterized in $\phi : \hat{\pi}(x; \phi)$
- Critic parameterized in $\theta : \hat{V}(x; \theta)$

Parameters ϕ and θ , have finite size, but approximate functions on continuous (infinitely large) spaces

2.2. Update Critic: Value Estimation

The task of the critic is to **predict the expected future reinforcement r the process will receive being in the current state and following the current control policy.**

- For doing that, we need to train critic, the prediction error is always a train input:

Use sample $(x_k; u_k; x_{k+1}; r_{k+1})$ at each step k and parameterized V :

Note that both $\hat{V}(s_k)$ and $\hat{V}(s_{k+1})$ are known at time k , since $\hat{V}(s_{k+1})$ is a prediction obtained for the current process state: For example, assuming a NN, we can always get a prediction of next state value based on current state, we get the Δ from **real system feedback** r_{k+1} and **critic original prediction**

$$\Delta_k = V(s_k) - \hat{V}(s_k) = r_{k+1} + \gamma \hat{V}(s_{k+1}) - \hat{V}(s_k) \quad (4)$$

- For example, Let the critic be represented by a neural network or a fuzzy system, in my opinion, **the index of θ represents the iteration version, it accidentally equal to state/step number, because each step we make a update. It does not mean the corresponding θ of state k or $k+1$**

$$\hat{V}(s_{k+1}) = \hat{V}(s_{k+1}, \theta_k)$$

To update θ_k , a gradient-descent learning rule is used:

$$\theta_{k+1} = \theta_k + \alpha_c \Delta_k \frac{\partial \hat{V}(s_k, \theta_k)}{\partial \theta_k} \quad (5)$$

$\alpha_c > 0$, learning rate of critic

$$\Delta_k > 0, \text{ i.e., } r_{k+1} + \gamma \hat{V}^\pi(x_{k+1}, \theta_k) > \hat{V}^\pi(x_k, \theta_k) \quad (6)$$

\Rightarrow old estimate too low, increase \hat{V}

$$\Delta_k < 0, \text{ i.e., } r_{k+1} + \gamma \hat{V}^\pi(x_{k+1}, \theta_k) < \hat{V}^\pi(x_k, \theta_k) \quad (7)$$

\Rightarrow old estimate too high, decrease \hat{V}

2.3. Update Actor: Policy Update

The actor (i.e., the policy) can be adapted in order to **establish an optimal mapping between the system states and the control actions.**

$$u_k = \hat{\pi}(x_k, \varphi_k) + \tilde{u}_k, \hat{\pi} = \text{actor}, \tilde{u}_k = \text{exploration} \quad (8)$$

$$\varphi_{k+1} = \varphi_k + \alpha_a \Delta_k \tilde{u}_k \frac{\partial \hat{\pi}(x, \varphi)}{\partial \varphi} \bigg|_{x=x_k} \quad (9)$$

$$\Delta_k > 0, \text{ i.e., } r_{k+1} + \gamma \hat{V}^\pi(x_{k+1}, \theta_k) > \hat{V}^\pi(x_k, \theta_k) \quad (10)$$

$\Rightarrow \tilde{u}_k$ had positive effect, move in that direction

$$\Delta_k < 0, \text{ i.e., } r_{k+1} + \gamma \hat{V}^\pi(x_{k+1}, \theta_k) < \hat{V}^\pi(x_k, \theta_k) \quad (11)$$

$\Rightarrow \tilde{u}_k$ had negative effect, move away from that direction

3. One Really Classical Example

Lecture Notes, Application 1