

The art of Measurement

1. Overview of Evaluation Techniques

1.1. Analytical Modeling

Positive

Negative

1.2. Simulation

Positive

Negative

1.3. Measurement

positive

Negative

2. Model-based Evaluation

2.1. How to design a model-based evaluation?

What is the goal

What to measure

How to measure

How to summarize and present the output?

2.2. An example

2.3. System Parameters (configuration)

2.4. Workload model

2.5. Factors

3. Metrics

4. Example of practical considerations when measuring time

4.1. Attention

4.2. Measuring a process's execution time

Example

Times:

Wall-Clock time

Multi-Core Situation

Drawbacks

4.3. Measuring time inside a program

C library: Gettimeofday & clock_gettime()

Using timers

4.4. Measuring time from outside of the system

Monitor

5. The Art of Data Gathering

5.1. Repeat Measurements

Errors

Uncertainties

5.2. Cold start

5.3. Other Considerations of Gathering Data

The effect of hidden or ignored parameters

Measurement-induced perturbations

Smart Compiler

6. The Art of Data Summarization

7. Common mistakes when design measurements

Biased Goals

Unsystematic Approach

Incorrect Performance Metrics

Unrepresentative Workload

Not performing a sensitivity analysis on the factor

Ignoring variability of results when reporting the outcome

Omitting Assumptions and limitations

1. Overview of Evaluation Techniques

1.1. Analytical Modeling

- Based on a rigorous mathematical models
- For example: Markov Chain, Queue Theory

Positive

- Provides the best insight into the effects of different parameters and their interaction
- Can be done before the system is built and takes a short time

Negative

- these models are rarely accurate
 - Depend on the quality and correctness of assumptions that made to simplify the model

1.2. Simulation

simulate the system operation

Positive

- **full control** of simulation model, parameters, level of detail (high flexibility)
- Can be done **before the system is built**

Negative

It may still include simplifying **assumptions**

1.3. Measurement

Implement the system in full and measure its performance directly

positive

The most **convincing**

Negative

has high costs

2. Model-based Evaluation

2.1. How to design a model-based evaluation?

What is the goal

Can you achieve what you want to achieve by measurement?

What to measure

- Metrics: Does your metric really measure what you want to measure
- Workload: What type of input workload should you consider

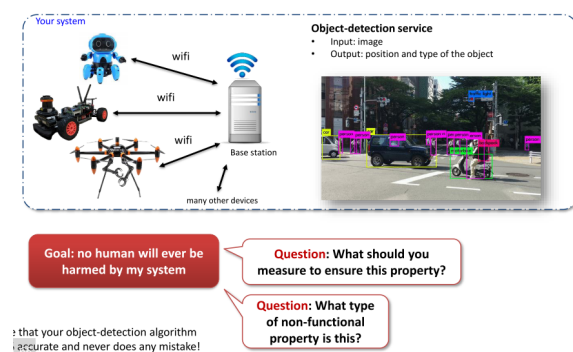
How to measure

- What **tools** should/can you use for measuring a metrics?
- **Considerations** while doing measurement

How to summarize and present the output?

- How to **summarize** your thousands of data points
- What are the **statistical methods/concepts** that can help you

2.2. An example



Metric types

- a **count** of how many times an event occurs
- the **duration** of some time interval
- the **size** of some parameter

Examples?

Metrics can have conditions

- Number of successfully processed images in one hour
- Total execution time of the image processing task for serving 100 images
- Time until the next failure


2.3. System Parameters (configuration)

System parameters are configurations and parameters that impact a system's performance **independent of the system's input workload**

Example:

Metric is the round trip time of the object tracking

Question: what are the system parameters in our example?




System parameters for our example:

- **Speed of:** base station's CPU, robot's microcontroller, and network.
- Scheduling policy on both ends.
- Network protocol.
- Operating system overhead for interfacing with the network.
- Reliability of the network affecting the number of retransmissions required.
- ...

2.4. Workload model

- Workload model defines the model of input workload to the system
- It is usually a representation of the **actual usage of the system** after deployment

Question: what are the workload parameters?



Workload parameters for this example:

- Number of IoT devices (user)
- Frequency of sending requests to the base station
 - and distribution (periodic, Poisson, ...)
- Complexity of the navigation task at different locations
 - Example: around cross roads is harder than straight roads
 - Proximity (presence) of other devices when it reaches to a cross road
- Other load on the robot's microcontroller and base station CPUs
- Other load on the network

2.5. Factors

- **Factors:** Parameters that are varied during the evaluation
- **Levels:** Values of a factor

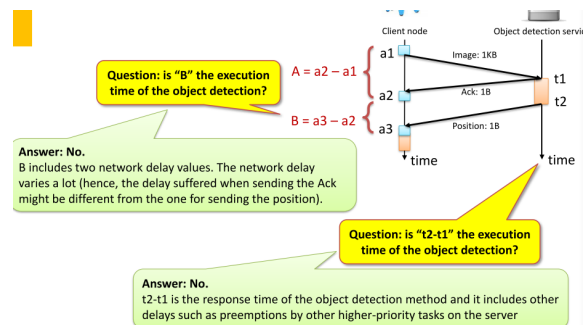
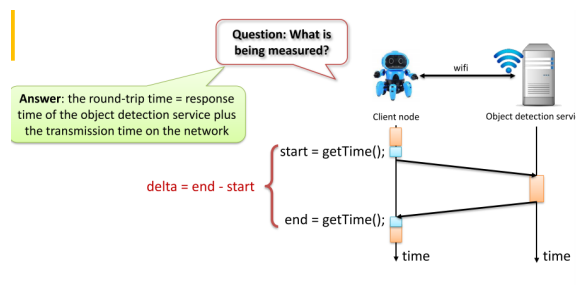
Example:

- Image size: {10KB, 100KB, 1MB}
- Frequency of calls to service:
 - Periodic: 1 every {1, 10, 100, 1000} milliseconds
 - Exponential distribution with rate: {0.1, 0.2, 0.3, 0.4, 0.5}

Assumption:

- Fixed: type of CPU and operating system.
- Measure under no other load on the hosts and the network.

3. Metrics



4. Example of practical considerations when measuring time

4.1. Attention

Don't measure by hand!

4.2. Measuring a process's execution time

We use UNIX's "time" command

Example

```
surf:~$ /usr/bin/X11/time ls -la -R ~/ > /dev/null
```

```
2.11user 1.94system 1:8.72elapsed 2%CPU
(0avgtext+0avgdata 0maxresident)k
0inputs+0outputs (0major+1155minor)pagefaults 0swaps
```

- 2.11 seconds of **user time**
- 1.94 seconds of **system time**
- 1 minutes and 8.72 seconds of **wall-clock time**
- 2% of CPU was used
- 0+0k memory used (text + data)
- 0 input, 0 output (file system I/O)
- 1155 minor pagefaults
- 0 swaps
- **User Time**
 - time spent executing user code
- **System Time**
 - time spent executing kernel code (inside OS)
- **Wall-Clock Time**
 - time from start to end

Times:

User time:

time spent executing **user code**

System time:

- time spent executing **kernel code** (inside OS)
- Larger system time may be a sign of higher calls to system APIs such as fopen(), fread(), fwrite(), etc

Wall-Clock time:

time **from start to end**

Wall-Clock time

Wall clock – system – user = I/O + preemption + suspension + other interferences

- **I/O:** time waiting for performing an I/O interaction (including reading from or writing to files, ports, etc.)
- **Preemption:** time waiting for other tasks to execute

- **Suspension:** time being blocked (e.g., when a task wants to access a data that is protected by a lock) or being suspended due to an I/O operation
- **Other interferences:** time spent on accessing memory bus of a memory bank, etc.

One way to control the impact of these delays is try to run your processs in isolation (when no other process is running)

Multi-Core Situation

```
surf:~$ /usr/bin/X11/time ./parallelQuicksort2
```

```
9.76user 10.51system 0:06.11elapsed 331%CPU
(0avgtext+0avgdata 158268maxresident)k
0inputs+0outputs (0major+7599minor)pagefaults 0swaps
```

Question: Why wall clock time < system + user time?

This system could have multiple cores and execute the program in parallel on multiple cores. In that case, the wall-clock time becomes smaller than user or system times.

Drawbacks

- The time command has poor resolution
 - “Only” milliseconds
- “time” times the **whole code**
 - Often ES code is not terminating (i.e. runs forever)
 - Sometimes we're only interested in timing some part of the code, for instance the one that we are trying to optimize

4.3. Measuring time inside a program

C library: Gettimeofday & clock_gettime()

- Measures the number of microseconds since midnight, Jan 1st 1970, expressed in seconds and microseconds
- Compute the time elapsed in microseconds: $(\text{end.tv sec}1000000.0 + \text{end.tv usec} \text{start.tv sec}1000000.0 - \text{start.tv usec}) / 1000000.0$

Using timers

Assumption: There is no other process or an operating system in this example. The goal is to measure the runtime of the middle instruction that runs on an Arduino Mega.



```
long x = 30, start, end;
start = micros();
x = 42 * x;
end = micros();
printf(end - start);
```

micros() returns values that are integer multiples of 4 microseconds

Each instruction is about 2~3 microsecond

What can invalidate this measurement?

- Timers are not infinitely accurate
 - All clocks have a granularity
 - The error in a time measurement, even if everything is perfect, may be the size of this granularity (sometimes called a clock tick)
- Always know your clock granularity
- Ensure that your measurement is for a long enough duration (say 100x the “tick”)



- high resolution
- early rollover

- low resolution
- late rollover

T'_c	32-bits	64-bits
10 ns	42 s	5058 years
1 μ s	1.2 hour	0.5 million years
1 ms	49 days	0.5×10^9

Time Rollover

- Occurs when an n-bit counter undergoes a transition from its maximum value 2^n-1 to zero.
- here is a **trade-off** between rollover time and granularity of the timer

4.4. Measuring time from outside of the system

Monitor

A Monitor is a tool to observe the activities on a system. In general, it

1. makes **measurements** on the system (Observation)
2. collects performance **statistics** (Collection),
3. **analyzes** the data (Analysis),
4. **displays** results (Presentation).

5. The Art of Data Gathering

5.1. Repeat Measurements

Errors

Error is the difference between **the measured value** and **the ‘true value’ of the thing being measured**

- Error can be caused by a **low timer resolution**

Uncertainties

Uncertainty is a **quantification** of the doubt about the measurement result

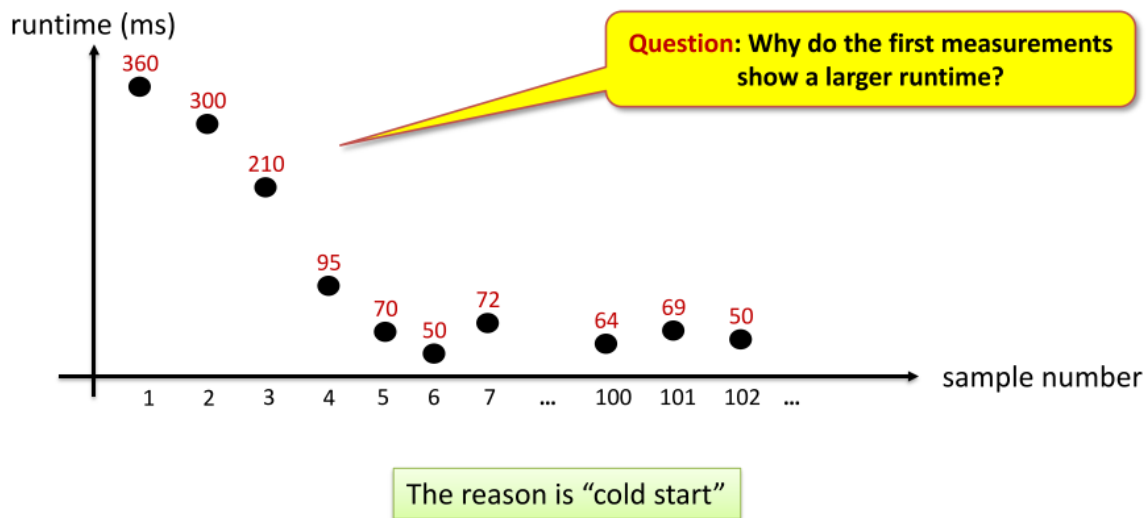
- Network delay and signal wifi strength can be two sources of uncertainty

5.2. Cold start

- Code may still be on disk, and not even loaded into memory.
- Data may be in slow memory rather than fast, e.g., cache

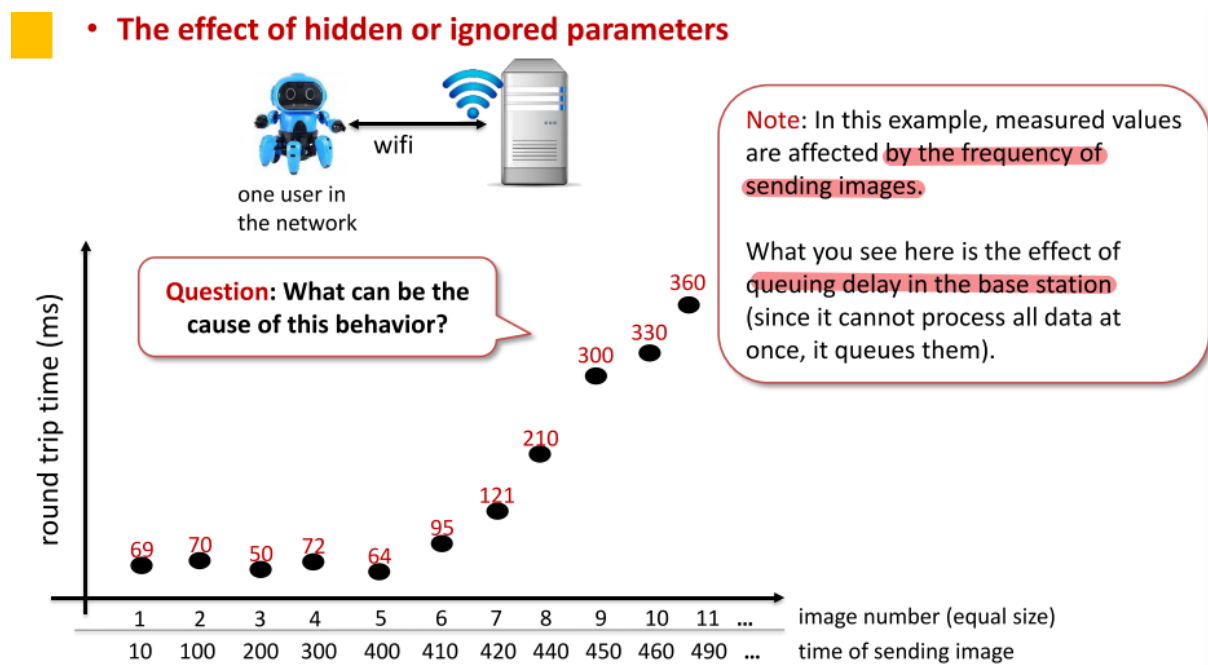
So, we should let the program run for a while and then start measuring

- usually you need to put effort to ensure that the data in the intended part of the memory hierarchy.



5.3. Other Considerations of Gathering Data

The effect of hidden or ignored parameters



Measurement-induced perturbations

The system resources consumed by the measurement tool itself as it collects data may strongly affect the system's performance.

On embedded systems, writing outputs on the terminal will heavily impact the system's behavior

In Arduino Mega, storing time in an array takes about 2~3 microseconds, while writing it using `printf()` is about 70~80 microseconds

- Do not write the measured data on the output while gathering samples
 - **Solution:** store the data in an array and then output the array after gathering samples
- Do not use dynamic memory allocation for storing samples
 - It **messes up the memory alignment** of the program and changes its behavior
 - You may even run into “**out of memory**” problem due to fragmentation of memory
 - **Solution:** Use an array with predefined size (static memory allocation)

Smart Compiler

If the result of the computation is not used, the compiler may eliminate the code.

- So sometimes, your measurement will be less than the actual time the program need

```
void bar(){
  long start = getTime();
  ...
  x = (y > 1) ? (y * y * 2) : (2.123456789 % 17);
  ...
  // foo(x);
  ....
  long end = getTime();
  print(end - start);
}
```

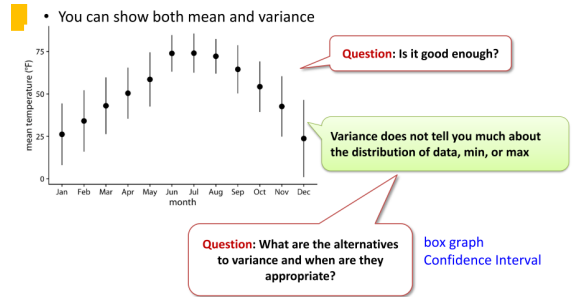
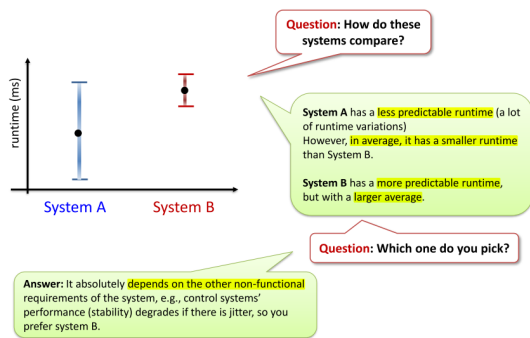
Question: When does it become an issue?

The compiler may remove the line that changes x, so your final measurement may look much shorter than it will be if “foo” was not commented out


You decided to comment out this line to avoid including the runtime of “foo” into “bar”.

6. The Art of Data Summarization

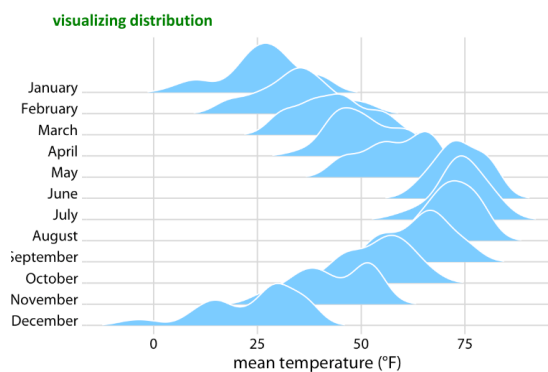

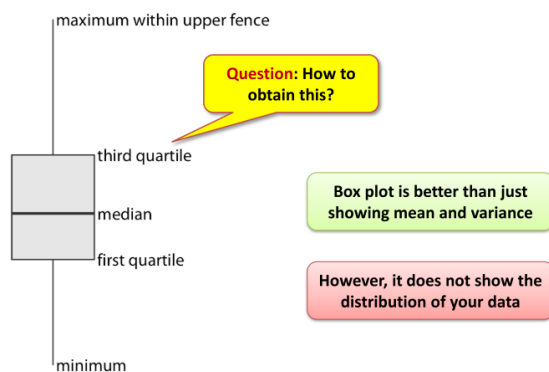
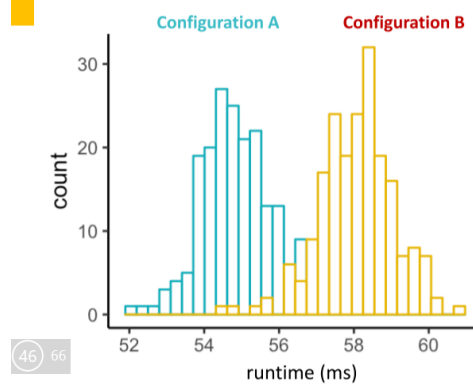
Interpreting results

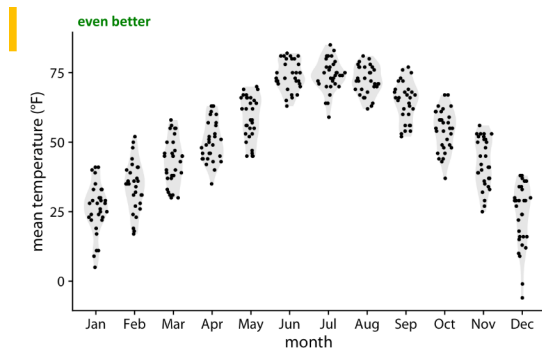


Confidence interval

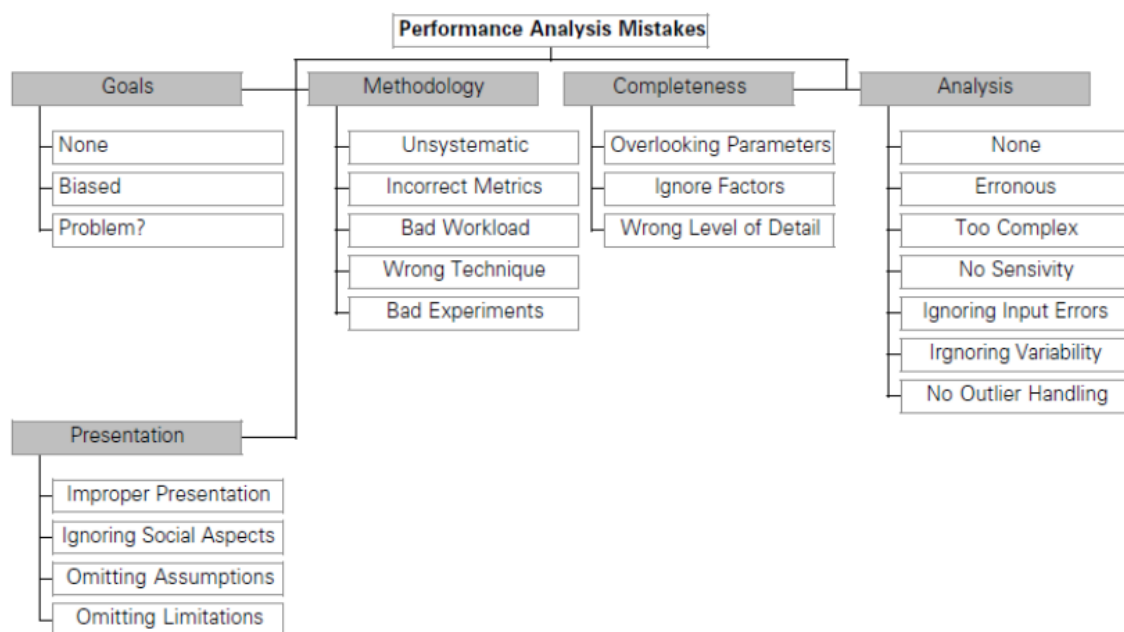


Histogram





7. Common mistakes when design measurements



Biased Goals

Example:

- To show that OUR system is better than THEIRS
- Mistake: Finding the metrics and workload such that OUR system turns out better rather than finding the right metrics and workloads

Unsystematic Approach

- Selecting system parameters , factors, metrics, and workload arbitrarily

Incorrect Performance Metrics

Example: Comparing Two CPUs based on the throughput (MIPS)

CISC

RISC

Unrepresentative Workload

- Example: if packet in the network are generally **mixture of long and short**, workload should consist of **short and long packet sizes**, not just long ones.

Not performing a sensitivity analysis on the factor

- Sensitivity analysis answers the question, "**if the factors deviate from the expectations, what will the effect be on the system being analyzed**".
- Putting too much emphasis on the results of the analysis, **presenting it as fact rather than evidence**.

Ignoring variability of results when reporting the outcome

- Only reporting mean value

Omitting Assumptions and limitations

- Assumption and limitations of the analysis are often omitted in the final report.
- This may lead the user to apply the analysis to another context where assumption will not be valid anymore.