

04_Aperiodic Task Scheduling

04_Aperiodic Task Scheduling

1. Earliest Due Date (EDD) Algorithm

- 1.1. Rules
- 1.2. Performance
 - Lateness
 - Jackson's Theorem
 - Proof:
- 1.3. Feasibility Check
- 1.4. Complexity

2. Earliest Deadline First (EDF) Algorithm

- 2.1. Rules
- 2.2. Performance (M.Dertouzos Theorem)
 - Proof
 - Property
- 2.3. Feasibility Check
- 2.4. Complexity

3. Non-preemptive Scheduling Algorithms

- 3.1. Model
- 3.2. Bratley's Algorithm
 - Pruning Rule
- 3.3. Spring Algorithm
 - Precedence Constraints

4. Scheduling with Precedence Constraints

- 4.1. Latest Deadline First (LDF) Scheduling ($1|Prec, Sync| L_{max}$)
 - Assumptions
 - Rules
 - Examples:
 - Property
 - Optimal Proof
- 4.2. EDF* Methods ($1|Prec, Preem| L_{max}$)
 - Assumptions
 - Basic Ideas
 - Rules
 - Property

Overview of Aperiodic Scheduling

Outline

- How to schedule aperiodic Tasks in 1 Processor?
 - Earliest Due Date (EDD) Algorithm
 - Definition, Examples, and Guarantee Tests
 - Earliest Deadline First (EDF) Algorithm
 - Definition, Examples, and Guarantee Tests
 - Non-preemptive Scheduling Algorithms
 - Bratley's Algorithm and Spring Algorithm
-
- Scheduling with Precedence Constraints
 - Latest Deadline First
 - Precedence Transformation and EDF

1. Earliest Due Date (EDD) Algorithm

1 | sync | L_{\max}

1.1. Rules

Rules:

select the task with the earliest **relative deadline**

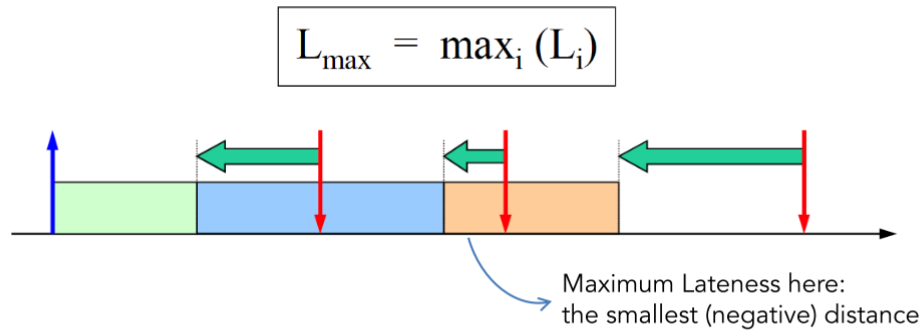
Assumptions:

- All tasks are activated **simultaneously** (like a "**batch**")
- Preemption is not used (since tasks arrive simultaneously);
- Static policy

1.2. Performance

Lateness

- If L_{max} is non-positive, then no task misses its deadline.
- **EDD minimizes L_{max}** (has the maximum slack time)



Jackson's Theorem

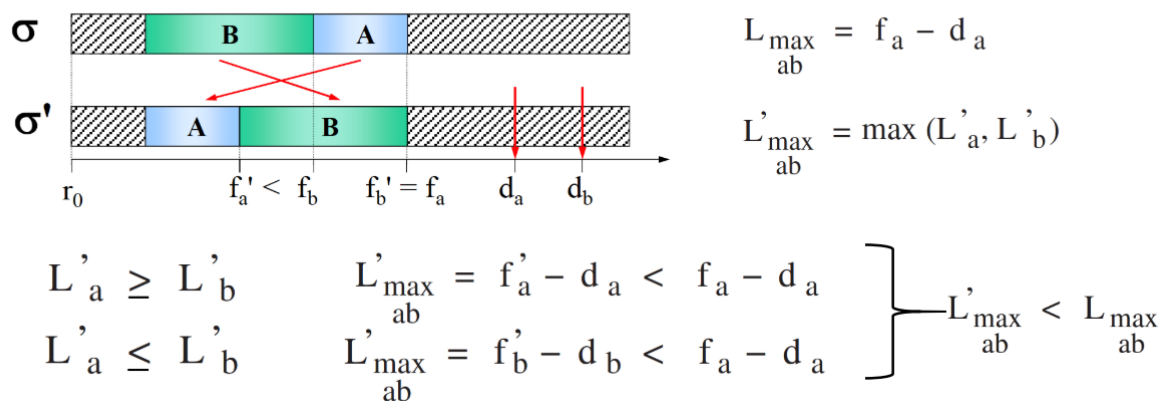
Given a set of n independent tasks, any algorithm that executes the tasks in order of increasing deadlines is **optimal w.r.t. minimizing L_{max}**

Proof:

Proof with two steps:

- optimal in one step

- In either case, we reduce the maximum lateness:



- recursively optimizing (like bubble sort)

If we continue doing these swaps, we will end up with the optimal, in terms of L_{\max} , EDD policy.

$$\sigma \rightarrow \sigma' \rightarrow \sigma'' \rightarrow \dots \rightarrow \sigma^*$$

$$L_{\max}(\sigma) \geq L_{\max}(\sigma') \geq L_{\max}(\sigma'') \dots \geq L_{\max}(\sigma^*)$$

1.3. Feasibility Check

Feasibility must be **checked offline**

- for each task should hold:

$$\forall i = 1, \dots, n \quad f_i \leq d_i$$

- Assume task indices are ordered in increasing deadlines, then, worst-case finishing time of each task is

$$f_i = \sum_{k=1}^i C_k$$

- feasibility requires to satisfy the following constraints

$$\forall i = 1, \dots, n \quad \sum_{k=1}^i C_k \leq d_i$$

1.4. Complexity

$$O(n \log(n))$$

2. Earliest Deadline First (EDF) Algorithm

2.1. Rules

execute the task with the earliest **absolute deadline**

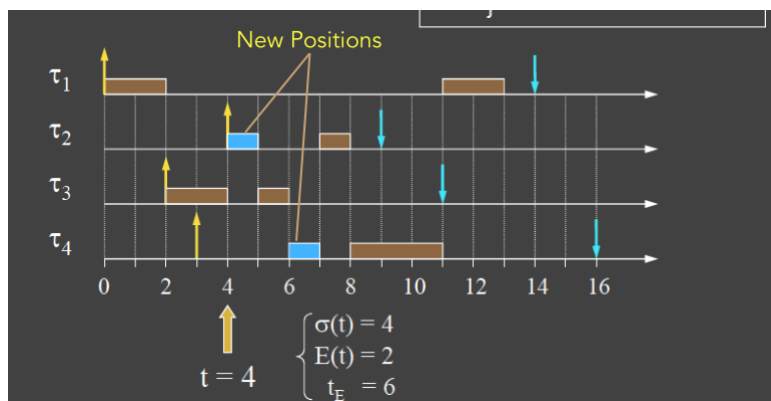
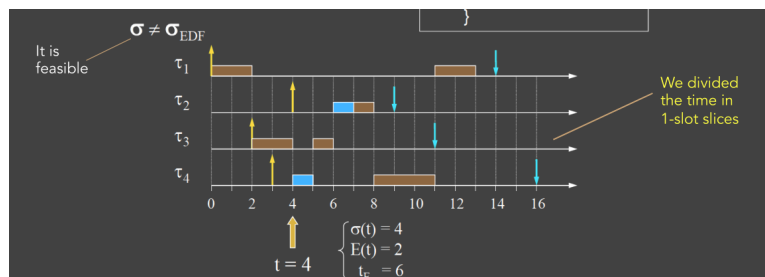
Assumption

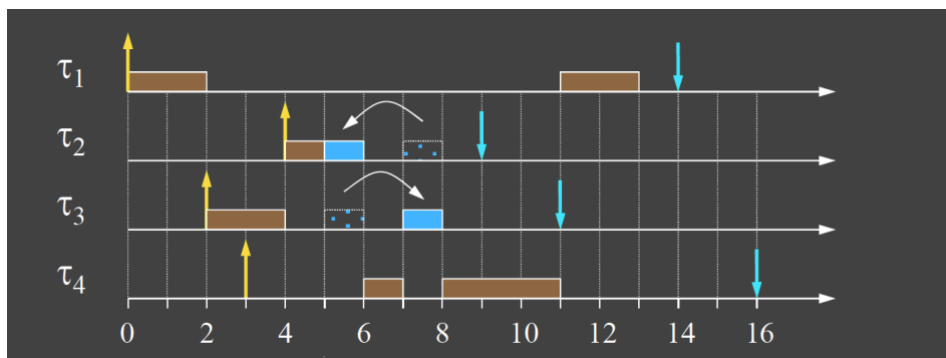
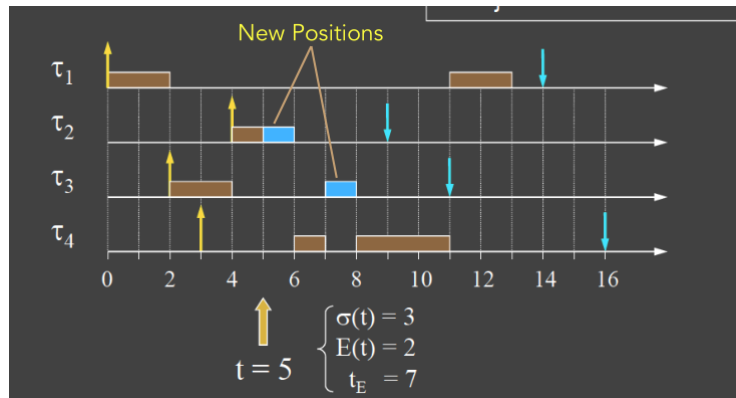
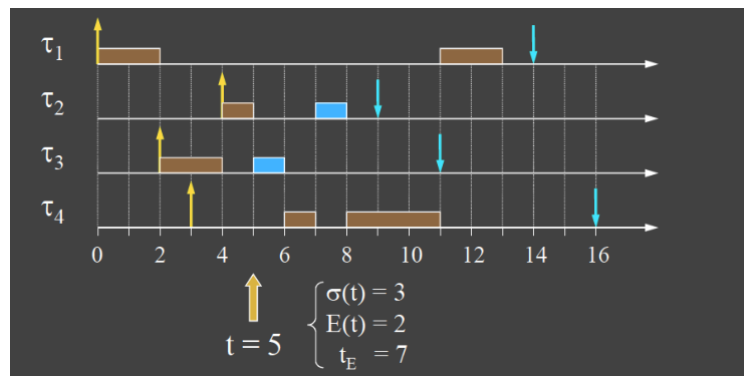
- tasks may arrive at any time (**asynchronous**);
- **Preemption** is used;
- **Dynamic** priority; (d_i depends on arrival)

2.2. Performance (M.Dertouzos Theorem)

Given a set of n independent tasks **with arbitrary arrival times**, any algorithm that at any instant executes the task with the earliest absolute deadline among all ready tasks **is optimal w.r.t. minimizing the maximum lateness** L_{\max}

Proof





Feasibility is also guarantee, and the lateness is improve

- If the later deadline one's finish time is not change \rightarrow absolutely hold
- If the later deadline one's finish time change (swap):
 - It can be easily prove the earlier deadline one is the L_{max} if the swap happen

Property

So, EDF is also optimal in the sense of feasibility. i.e. **If there exists a feasible schedule for this set of tasks, EDF will find it**

2.3. Feasibility Check

- The test runs upon arrival of every new task
- we perform **admission control** by checking if the new task set is schedulable
- check if **under worst-case scenario** the tasks can meet their deadlines.
- Tasks are preemptable, with C_i being the initial computation time and $c_i(t)$ the remaining comp time at t

$$f_i = \sum_{k=1}^i c_k(t)$$

$$\forall i = 1, \dots, n \quad \sum_{k=1}^i c_k(t) \leq d_i$$

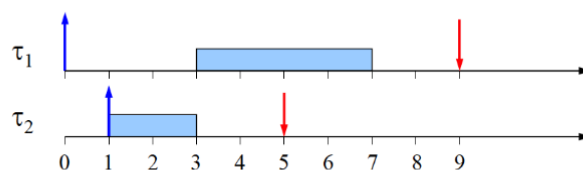
2.4. Complexity

$$O(n)$$

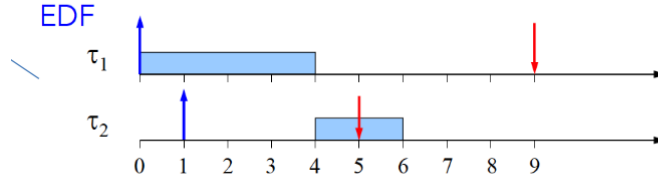
3. Non-preemptive Scheduling Algorithms

- EDF is **no longer optimal** under no preemption (with idle choice)
 - Results can be improved if we know the future, we can first let the CPU be idle for 1 slot until task 2 come
- EDF is still optimal if we focus on **non-idle** algorithms
 - Two algorithms in this section consider **with-idle** situation

Feasible schedule

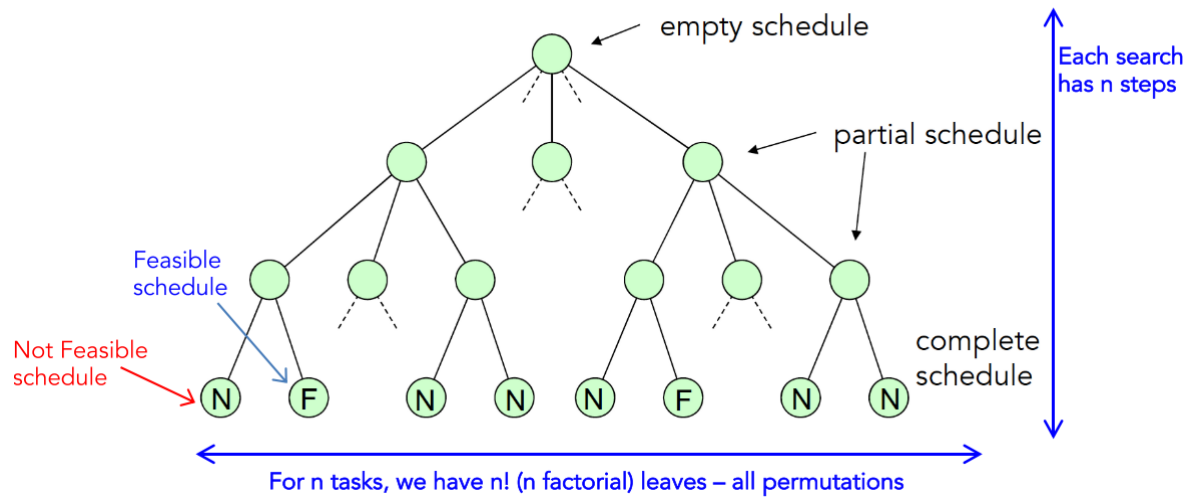


EDF



3.1. Model

Even if task arrivals are known, we must **search a tree** with $n!$ leaves and n depth; hence $O(nn!)$ worst-case complexity



3.2. Bratley's Algorithm

$$1|no-preem|L_{max}$$

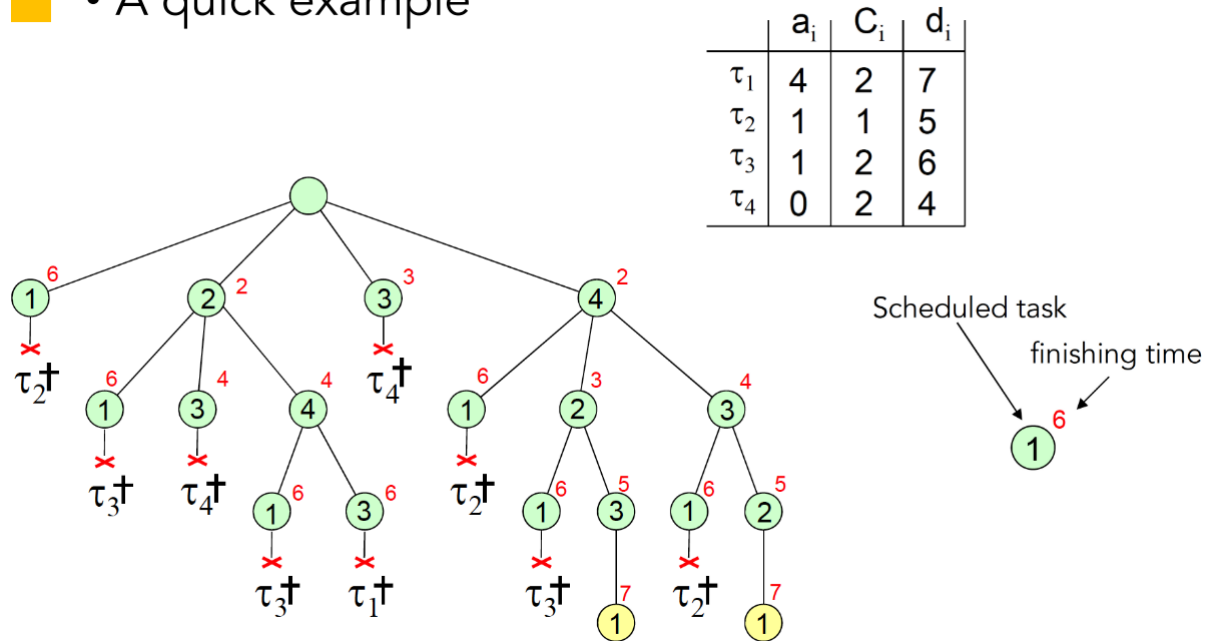
Pruning Rule

Stop searching a path:

- If a feasible schedule is **already found**; or
- If the addition of **any node** to current path causes a missed deadline



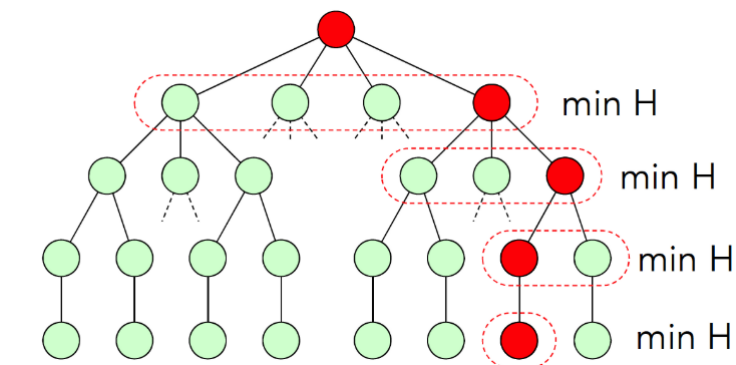
• A quick example



3.3. Spring Algorithm

Reduce complexity by sacrificing optimality: **Heuristics**

At each step select the task **minimizing a heuristic function H** (n steps)



- Optimizing one criterion:

$$H = r_i \Rightarrow \text{FCFS}$$

$$H = C_i \Rightarrow \text{SJF}$$

$$H = d_i \Rightarrow \text{EDF}$$

- Optimizing multiple criteria:
 - **Pareto** optimization through scalarization

$$H = w_1 r_i + w_2 D_i$$

$$H = w_1 C_i + w_2 d_i$$

$$H = w_1 V_i + w_2 d_i$$

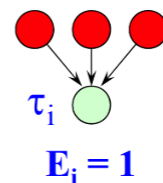
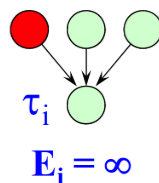
Precedence Constraints

Use "Flag" parameter

$$H = E_i (w_1 r_i + w_2 D_i)$$

$$H = E_i (w_1 C_i + w_2 d_i)$$

Task i can be executed
only if these three
Tasks finish



4. Scheduling with Precedence Constraints

4.1. Latest Deadline First (LDF) Scheduling (1|Prec,Sync| L_{max})

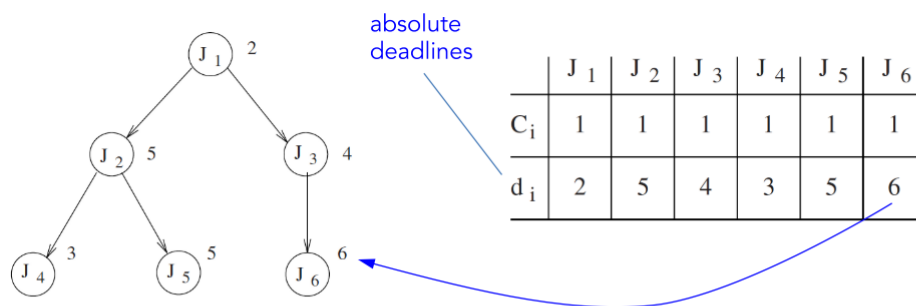
Assumptions

Synchronous Activation of Tasks

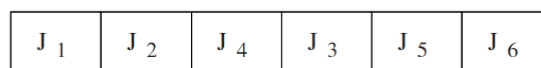
Rules

- Take as **input** the set of tasks and their **precedence graph (DAG)**;
- Add task from the end of the queue:
 - Construct the schedule starting **from the tail of the DAG**;
 - **Among tasks with no successors (or, already-selected successors), pick the task with the latest deadline to execute last.**
 - At runtime, select first tasks from the head of queue.

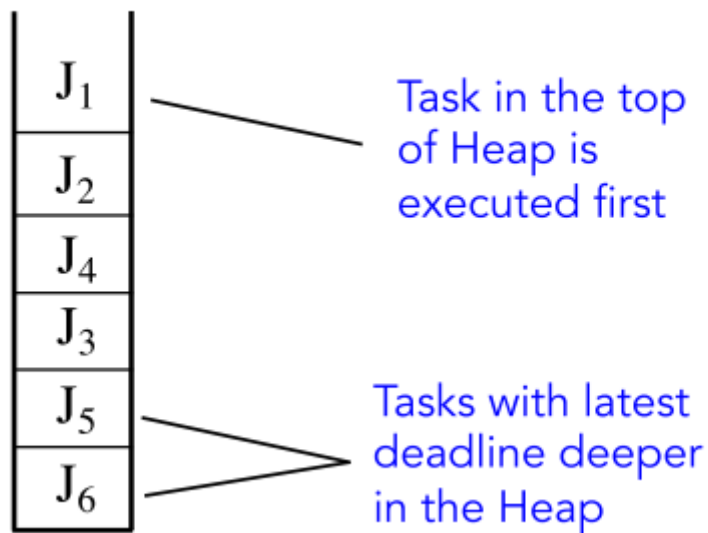
Examples:



- Start from the end of queue; select tasks with no "children" or already-selected children.

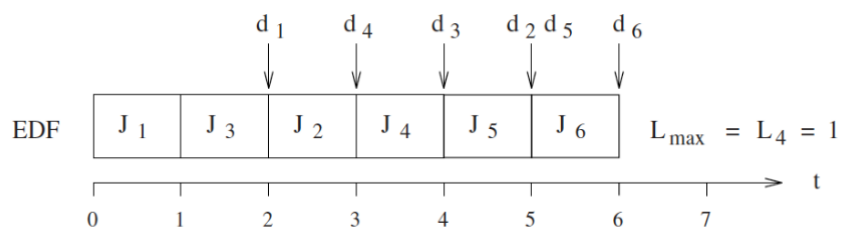


We begin at the end from end build the queue



Notice:

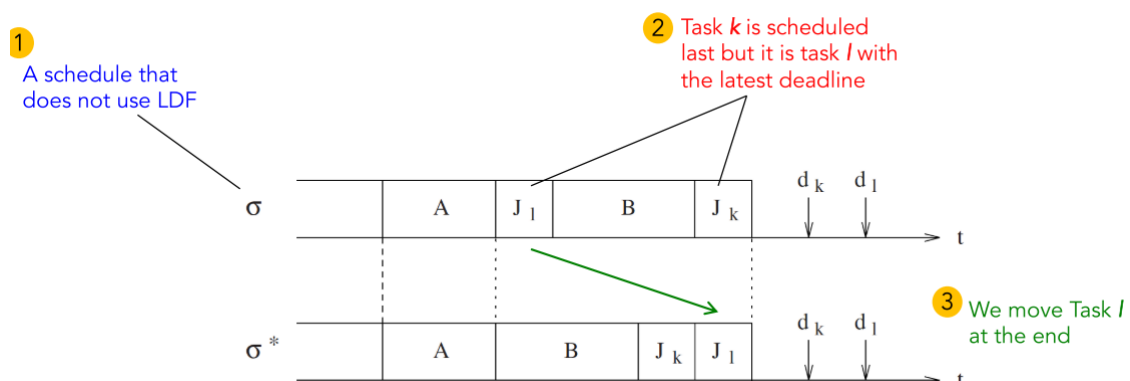
EDF would have produced an **infeasible** schedule: select the task with the **earliest deadline** among all **currently-eligible tasks**



Property

- polynomial on number of tasks n
- **Optimal in terms of minimum L_{\max}**

Optimal Proof



L_{max} of the new schedule is reduced:

- For Tasks in set A, we do not have any change in $L = \{\text{finish} - \text{deadline}\}$ times;
- For Tasks in set B, we have, actually, reduced (improve) the Lateness;
- For Task k we have reduced the Lateness;
- For Task l the lateness is no worse than the L_{max} of schedule σ since $f - d_l < f - d_k$

So for three of the part, they have better L, for one part, it has better L than one part of the original schedule

$$\begin{aligned} a &> a'; b > b'; c > c'; c > d'; \\ \max(a, b, c, d) &\geq c; \\ \max(a', b', c', d') &< c \end{aligned}$$

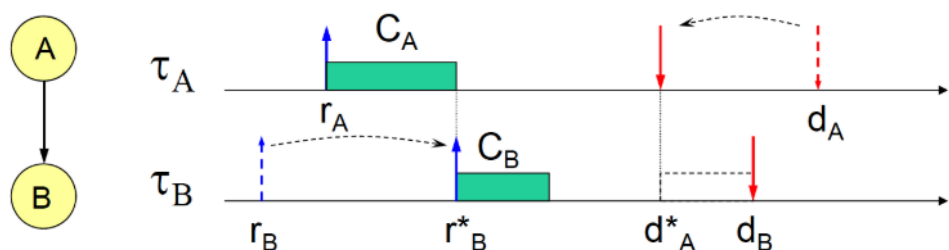
4.2. EDF* Methods (1|Prec,Preem| L_{max})

Assumptions

Arrival times must be known in advance

Basic Ideas

- Transform task set Γ to an equivalent set Γ^* with no Prec. constraints;
 - Postpone the release time of a successor;
 - Advance the deadline of a predecessor
 - E.G. Task A is a predecessor and in the new schedule Γ^* should finish:



- Before its original deadline: d_A
 - Before the maximum start time of Task B: $d_B - C_B$
- Apply EDF to the new task set Γ^*

Rules

- Arrival Time Modification: (downward)
 - a. For all root nodes, set: $r_i^* = r_i$
 - b. Select a task i that:
 - its release time has not been modified;
 - the release times of all its immediate predecessors have been modified;
 - c. Set: $r_i^* = \max \{r_i, \max_{\tau_k \rightarrow \tau_i} (r_k^* + C_k)\}$
 - d. GoTo b.
- Deadline Modification: (upward)
 - a. For all Leaf nodes, set: $d_i^* = d_i$
 - b. Select a task i that:
 - its deadline has not been modified;
 - the deadlines of all its immediate successors have been modified
 - c. Set: $d_i^* = \min \{d_i, \min_{\tau_i \rightarrow \tau_k} (d_k^* - C_k)\}$
 - d. GoTo b.

Property

- If Γ^* is schedulable with EDF, then Γ is also schedulable with EDF
 - because in Γ^* , we have $r_i^* \geq r_i$ $d_i^* \leq d_i$
- The precedence constraints of Γ are satisfied in Γ^*
 - if $\tau_i \rightarrow \tau_j$ then $d_i^* < d_j^*$ and $r_i^* < r_j^*$

Overview of Aperiodic Scheduling

	sync. activation	preemptive async. activation	non-preemptive async. activation
independent	EDD (Jackson '55) $O(n \log n)$ Optimal	EDF (Horn '74) $O(n^2)$ Optimal	Tree search (Bratley '71) $O(n n!)$ Optimal
precedence constraints	LDF (Lawler '73) $O(n^2)$ Optimal	EDF * (Chetto et al. '90) $O(n^2)$ Optimal	Spring (Stankovic & Ramamritham '87) $O(n^2)$ Heuristic